

# Hyogi Sim

Oak Ridge National Laboratory  
simh@ornl.gov, +1-865-574-6167

## Current Position

---

HPC Systems Engineer

*National Center for Computational Sciences, Oak Ridge National Laboratory*

## Education

---

Ph.D. in Computer Science, Virginia Tech, VA

— *expected in 2020*

- Advisor: Dr. Ali R. Butt
- Thesis: Exploiting Storage-Side Computing Power for Fast and Efficient Scientific Data Discovery (*working title*)

M.S. in Computer Science, Virginia Tech, VA

— *Dec. 2014*

- Advisor: Dr. Ali R. Butt
- Thesis: AnalyzeThis: An Analysis Workflow-Aware Storage System

M.S. in Electronics and Computer Engineering, Hanyang University, S. Korea

— *Feb. 2008*

- Thesis: A Study of Performance Impact of Merging Storage Layers on Flash-Based DBMS

B.S. in Urban Planning, Hanyang University, S. Korea

— *Feb. 2005*

## Work History

---

HPC Systems Engineer, Oak Ridge National Laboratory

*Feb. 2017 – present*

Postmasters Research Assistant, Oak Ridge National Laboratory

*Mar. 2015 – Jan. 2017*

Summer Intern, Oak Ridge National Laboratory

*Jan. 2013 – Aug. 2013*

Graduate Research and Teaching Assistant, Virginia Tech

*Aug. 2011 – Dec. 2013, Sep. 2013 – Feb. 2015*

Senior Developer, Pitapat Mobile, S. Korea

*Jan. 2011 – June 2011*

Part-Time Programmer, Dept. of English Education, Hanyang University, S. Korea

*Aug. 2009, Feb. 2010*

Part-Time Programmer, Metabuild co., S. Korea

*Dec. 2008*

Senior Developer, Golfschool co., S. Korea

*2002 – 2003*

Military Conscript in Combat Police Force, S. Korea

*May 1999 - Jul. 2001*

## Skills & Expertise

---

Skillful in Linux application/system/kernel programming, parallel/distributed file systems, database systems.

- Programming Languages: C, C++, Python, JAVA, PHP, C#, BASH
- File & Storage: FUSE, Linux VFS, Device Mapper, SCSI OSD-2, Linux TGT, blktrace, exofs, ext3
- Parallel & Distributed File Systems: GlusterFS, Lustre, Ceph, NFS, HDFS
- Non-Volatile Memory & SSD: Linux MTD, Flash Translation Layer
- Databases: MySQL, PostgreSQL, SQLite, MS SQLServer, IBM DB2, WiSS, HyperDex
- Parallel Programming: MPI, Pthread, OpenMP, Spark
- Parallel I/O Library: netCDF, HDF5, SCR
- I/O Benchmarks: fio, mdtest, IOR, TPC-C, TPC-H
- Tools: gcc, gdb, cscope, ctags, autotools, git, svn, eclipse, glade, visual studio, latex, gnuplot
- Web Development: HTML, CSS, PHP, Javascript, JQuery, ASP.NET
- Package Management: RPM, Portage (Gentoo Linux), Pkgsrc (NetBSD)

## Project Experience in File and Storage Systems

---

**A metadata indexing framework for HPC archival storage systems**

*2018 – present, Oak Ridge National Lab.*

- Developing a software framework that periodically collects metadata from HPSS and builds data catalog.
- Developing a framework that extracts metadata from well-known document files without user intervention.
- Developing command-line utilities for system administrators and users.
- (*HPSS, IBM DB2, PostgreSQL, Apache Tika, XDR, cURL*)

**Development of a userspace checkpoint file system for HPC burst buffers**

*2017 – present, Oak Ridge National Lab.*

- Designing and developing an ephemeral, distributed file system for node-local burst buffers to facilitate checkpointing of HPC applications.

- (*FUSE, MPI, MDHIM, LevelDB, Mercury, Margo, Argobots*)

## **Large-scale file system snapshot and workload analysis**

2016 – 2019, Oak Ridge National Lab.

- Analyzing the snapshot data of Spider II (32 PB Lustre) file system.
- Analyzing the eight year worth workloads of 80 PB HPSS archival storage system.
- (*Lustre, HPSS, Spark, SparSQL, Parquet, Python, MySQL, SQLite*)

## **File system-integrated search and discovery services for HPC**

2015-2017, Oak Ridge National Lab.

- Designing and developing a file system-integrated metadata indexing framework that supports user-defined tagging in GlusterFS and CephFS.
- (*GlusterFS, CephFS, Linux, SQLite*)

## **Analysis-aware storage system for high performance computing**

2013 – 2015, Oak Ridge National Lab./Virginia Tech

- Designing an active execution framework based on SCSI T10 OSD-2 specification and implemented the extended OSD-2 protocol on the Linux TGT.
- Extending the Linux exofs file system to support the extended OSD protocol.
- Developing a FUSE file system that manages the array of active OSD devices.
- (*Linux Kernel, SCSI T10 OSD-2 Protocol, Linux TGT, FUSE, SQLite*)

## **Hierarchical data management in media servers with hybrid storage architecture**

2007 – 2009, Hanyang University

- Developing a content popularity analyzer that dynamically analyzes incoming requests to determine popular media contents in a commercial media server.
- Developing a light-weight file system for storage class memory devices (PRAM) in media cache servers.
- (*Linux Kernel, Windows Media Server*)

## **Development of a NAND Flash Memory-Based File System Supporting Transaction and Record Structure**

2006 –

2008, Hanyang University

- Developing a DBMS that directly manages a NAND flash memory via Linux MTD layer for storing EPG (Electronic Program Guide) data in a set-top box.
- Developing a new PostgreSQL buffer manager that generates an optimized LBA sequence for NAND flash memory.
- Developing a framework that generates a set of pre-defined I/O requests, measures latency, and identifies a FTL mapping algorithm of a NAND flash memory-based storage device.
- (*Linux Kernel, Linux MTD, Wisconsin Storage System, PostgreSQL, ARM-based embedded board*)

## **Other Project Experience**

### **A programming framework for processing-in-memory architecture**

2015 – 2016, Oak Ridge National Lab.

- Developing a high-level programming interface (C/C++) and a runtime environment that facilitates data analysis tasks with an array of PIM devices.
- Developing a kernel device driver that emulates PIM devices in a NUMA architecture.
- (*NUMA, Pthread, Intel TBB, Linux Kernel*)

### **Development of a web-based social marketing game**

2011, Pitapat Mobile, S. Korea

- Developing a social quiz/marketing web application using the Microsoft ASP.NET.
- Integrating the web application to the Facebook app platform using the Facebook API.
- (*C#, Javascript, JQuery, HTML, PHP, SQL Server, MySQL, Amazon EC2*)

### **Development of an image transfer library for a self-driving vehicle**

2008, Metabuild co., S. Korea

- Developing a userspace library that transfers requested pixel data from the CMOS camera to a lane-recognition module in a self-driving vehicle.
- (*Linux, C, ARM-based embedded board*)

### **Development of an online survey website**

Feb 2010, Aug 2009, Dept. of English Education, Hanyang University, S. Korea

- Developing an online-survey web application including administrative tools.
- Developing a web application that visualizes interactions among educators.
- (*Linux, PHP, HTML, Javascript, MySQL, GD, Apache, Microsoft Excel*)

### **Construction of a protein function database**

2006 – 2008, Hanyang University

- Developed a GUI environment that allows medical researchers to verify protein interaction graphs that are automatically extracted from academic articles.
- (*Windows, SQLServer, Delphi*)

### **Development of an online lesson website**

2002 – 2003, Golfschool co., S. Korea

- Developed and maintained a commercial company website that serves online golf lessons and an online shopping mall of sporting goods.
- (*Linux, Apache, PHP, MySQL, HTML, Javascript*)

## Publications

---

- Hyogi Sim, Sudharshan S. Vazhkudai, *Profiling the Usage of an Extreme-Scale Archival Storage System*, Proceedings of the 27th IEEE International Symposium on Modeling, Analysis and Simulation of Computer and Telecommunication Systems (MASCOTS '19), Rennes, France, October 2019
- Hyogi Sim, Arnab K. Paul, Eli Tilevich, Ali R. Butt, *CSLIM: Automated Extraction of IoT Functionalities from Legacy C Codebases*, Proceedings of the 8th International Workshop on Computing and Networking for IoT and Beyond (ComNet-IoT '19) in conjunction with ICDCN '19, Bangalore, India, pages 6, January 2019
- Ali Anwar, Yue Cheng, Hai Huang, Dongyoon Lee, Jingoo Han, Hyogi Sim, Fred Douglass, and Ali R. Butt, *BESPOKV: Application Tailored Scale-Out Key-Value Stores*, Proceedings of the 2018 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis (SC '18), Dallas, TX
- Hyogi Sim, Geoffroy R. Vallée, Youngjae Kim, Sudharshan S. Vazhkudai, Devesh Tiwari, Ali R. Butt, *An Analysis Workflow-Aware Storage System for Multi-Core Active Flash Arrays*, IEEE Transactions on Parallel and Distributed Systems (TPDS), Vol. 29, 2018
- Sangkuen Lee, Hyogi Sim, Sudharshan S. Vazhkudai, *A Programmable Shared-Memory System for an Array of Processing-In-Memory Devices*, Cluster Computing: The Journal of Networks, Software Tools and Applications, 2018
- Feiyi Wang, Hyogi Sim, Cameron Harr, Sarp Oral, *Diving into Petascale Production File Systems through Large Scale Profiling and Analysis*, In Proceedings of the 2nd Joint International Workshop on Parallel Data Storage & Data Intensive Scalable Computing Systems (PDSW-DISC '17) in conjunction with SC, Denver, CO
- Hyogi Sim, Youngjae Kim, Sudharshan S. Vazhkudai, Geoffroy R. Vallée, Seung-Hwan Lim, Ali R. Butt, *TagIt: An Integrated Indexing and Search Service for File Systems*, Proceedings of the 2017 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis (SC '17), Denver, CO
- Seung-Hwan Lim, Hyogi Sim, Raghul Gunasekaran, Sudharshan S. Vazhkudai, *Scientific User Behavior and Data-Sharing Trends in a Petascale File System*, Proceedings of the 2017 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis (SC '17), Denver, CO
- Sangkuen Lee, Hyogi Sim, Youngjae Kim, Sudharshan S. Vazhkudai, *AnalyzeThat: A Programmable Shared-Memory System for an Array of Processing-In-Memory Devices*, IEEE/ACM International Symposium on Cluster, Cloud and Grid Computing (CCGRID '17), Madrid, Spain, May 2017
- Hyogi Sim, Youngjae Kim, Sudharshan S. Vazhkudai, Geoffroy R. Vallée, Seung-Hwan Lim, Ali R. Butt, *TagIt: An Integrated Search and Discovery Service for Extreme-Scale File Systems*, Poster in the 2016 USENIX Annual Technical Conference (ATC '16), Denver, CO, June 2016
- Hyogi Sim, Youngjae Kim, Sudharshan S. Vazhkudai, Devesh Tiwari, Ali Anwar, Ali R. Butt, Lavanya Ramakrishnan, *AnalyzeThis: An Analysis Workflow-Aware Storage System*, Proceedings of the 2015 ACM/IEEE International Conference for High Performance Computing, Networking, Storage and Analysis (SC '15), Austin, TX
- Hyogi Sim, Youngjae Kim, Sudharshan S. Vazhkudai, Devesh Tiwari, Ali Anwar, Ali R. Butt, Lavanya Ramakrishnan, *AnalyzeThis: An Analysis Workflow-Aware Storage System*, Poster in the 2015 USENIX Annual Technical Conference (ATC '15), Santa Clara, CA, July 2015
- Hyogi Sim, Hoyoung Jung, Sungmin Park, Sooyong Kang, Jaehyuk Cha, *Identifying the FTL Mapping Scheme for USB Flash Devices*, The 4th International Conference on Convergence Technology and Information Convergence, CTIC 2009, Oct. 12
- Sooyong Kang, Sungmin Park, Hoyoung Jung, Hyogi Sim, Jaehyuk Cha, *Performance Tradeoffs in Using NVRAM Write Buffer for Flash Memory-based Storage Devices*, IEEE Transactions on Computers, Vol. 58, Issue 6 (Jun. 2009) Pages 744-758
- Hoyoung Jung, Hyogi Sim, Sungmin Park, Sooyong Kang, Jaehyuk Cha, *LRU-WSR: Integration of LRU and Writes Sequence Reordering for Flash Memory*, IEEE Transactions on Consumer Electronics, Volume 54, Issue 3 (Aug. 2008)
- Sungmin Park, Hoyoung Jung, Hyogi Sim, Sooyong Kang, Jaehyuk Cha, *Using Non-Volatile RAM as a Write Buffer for NAND Flash Memory-based Storage Devices*, 2008 IEEE International Symposium on Modeling, Analysis & Simulation of Computer & Telecommunication Systems, MASCOTS 2008, Sept 8-10, Baltimore, MD
- Sungmin Park, Hoyoung Jung, Hyogi Sim, Sooyong Kang, Jaehyuk Cha, *Write Buffer-aware Address Mapping for NAND Flash Memory Devices*, 2008 IEEE International Symposium on Modeling, Analysis & Simulation of Computer & Telecommunication Systems, MASCOTS 2008, Sept 8-10, Baltimore, MD
- Hoyoung Jung, Kyunghoon Yoon, Hyogi Sim, Sungmin Park, Sooyong Kang, Jaehyuk Cha, *LIRS-WSR: Integration of*

## **Presentations**

---

- Hyogi Sim, *Making a Peta-Scale Archival Storage System Searchable*, High Performance Storage Systems User Forum 2019 (HUF 2019), Indiana University, Bloomington, IN, October 2019
- Hyogi Sim, *Profiling the Usage of an Extreme-Scale Archival Storage System*, High Performance Storage Systems User Forum 2019 (HUF 2019), Indiana University, Bloomington, IN, October 2019

## **Recognitions**

---

- Analysis Restaurant: In November 2015, *DEIXIS online magazine* featured AnalyzeThis storage system as a monthly highlight, based on the SC '15 paper (<https://deixismagazine.org/2015/11/analysis-restaurant/>).

## **Graduate-Level Courses**

---

- Virginia Tech, VA: Statistics in Research, Multiprocessor Programming, Research Method in Computer Science, Advanced Parallel Computation, Advanced Topics in System and Network Security, Operating Systems, Software Refactoring
- Hanyang University, S. Korea: Advanced Operating System, Computer Algorithms, Database System Implementation, Real-Time Systems, Cryptography, Database Tuning

## **Teaching Experience**

---

Computer Organization II, Virginia Tech, VA	— Spring 2012, Fall 2011 (Teaching assistant)
File Structure, Hanyang University, S. Korea	— Spring 2009, Fall 2007 (Leading lab sessions)
Data Structures, Hanyang University, S. Korea	— Spring 2007 (Leading lab sessions)
Object Oriented Programming, Hanyang University, S. Korea	— Spring 2005 (Leading lab sessions)

## **Awards & Scholarships**

---

Graduate Research Assistantship, Virginia Tech, VA	— Fall 2011, 2012, 2013, Spring 2012, 2014
BK21 Scholarship, Hanyang Univeristy, S. Korea	— 2009
Academic Record Scholarship, Hanyang Univeristy, S. Korea	— 2005